

Práctica Hive - Spark

1. Crear las tablas payments, passenger, distance, tolls y congestion en Hive:

```
CREATE TABLE payments (VendorID int, tpep_pickup_datetime date, payment_type int,
total_amount float)
COMMENT 'Payments'
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ',';

CREATE TABLE tolls (tpep_pickup_datetime date, passenger_count int, tolls_amount float,
total_amount float)
COMMENT 'Tolls'
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ',';

CREATE TABLE passenger (tpep_pickup_datetime date, passenger_count int, total_amount float)
COMMENT 'Passenger'
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ',';

CREATE TABLE congestion (tpep_pickup_datetime date, passenger_count int,
congestion_surcharge float, total_amount float )
COMMENT 'Congestion'
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ',';

CREATE TABLE distance (tpep_pickup_datetime date, passenger_count int, trip_distance float,
total_amount float )
COMMENT 'Distance'
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ',';
```

Para ver los resultados en la terminal se utiliza *use tripdb* y *show tables*.

```
hive> use tripdb;
OK
Time taken: 0.05 seconds
hive> show tables;
OK
congestion
distance
passenger
payments
tolls
Time taken: 0.089 seconds, Fetched: 5 row(s)
```

2. En Hive hacer un ‘describe’ de las tablas passengers y distance:

```
hive> DESCRIBE passenger;
OK
tpep_pickup_datetime    date
passenger_count         int
total_amount            float
Time taken: 0.096 seconds, Fetched: 3 row(s)
```

```
hive> DESCRIBE distance;
OK
tpep_pickup_datetime    date
passenger_count         int
trip_distance           float
total_amount            float
Time taken: 0.118 seconds, Fetched: 4 row(s)
```

3. Hacer ingest del file 'Yellow_tripdata_2021-01.csv'.

```
hadoop@401bec58e4c6:~/landing$ hdfs dfs -ls /ingest
Found 2 items
-rw-r--r--    1 hadoop supergroup      5462 2024-04-19 23:34 /ingest/starwars.csv
-rw-r--r--    1 hadoop supergroup 125981363 2024-05-07 10:36 /ingest/yellow_tripdata_2021-01.csv
```

Para los siguientes ejercicios se debe usar Pyspark. Si se desea practicar más también se puede repetir los mismos en SQL.

4. Inserción de tablas mediante Pyspark con dataframes:

Carga del csv alojado en hdfs en el dataframe df.

```
df = spark.read.option("header", "true").csv("/ingest/yellow_tripdata_2021-01.csv")
df.show(1)
```

Para la tabla payments:

```
df_payments = df.select(df VendorID.cast("int"), df.tpep_pickup_datetime.cast("date"),
df.payment_type.cast("int"), df.total_amount.cast("float"))
df_payments.write.insertInto("tripdb.payments")
```

Para la tabla passenger:

```
df_passenger = df.select(df.tpep_pickup_datetime.cast("date"), df.passenger_count.cast("int"),
df.total_amout.cast("float"))
df_passenger.write.insertInto("tripdb.passenger")
```

Para la tabla tolls

```
df_tolls = df.select(df.tpep_pickup_datetime.cast("date"), df.passenger_count.cast("int"),
df.tolls_amount.cast("float"), df.total_amount.cast("float"))
df_tolls.write.insertInto("tripdb.tolls")
```

Para la tabla congestion:

```
df_congestion = df.select(df.tpep_pickup_datetime.cast("date"), df.passenger_count.cast("int"),
df.congestion_surcharge.cast("float"), df.total_amount.cast("float"))
df_congestion.write.insertInto("tripdb.congestion")
```

Para la tabla distance:

```
df_distance = df.select(df.tpep_pickup_datetime.cast("date"), df.passenger_count.cast("int"),
df.trip_distance.cast("float"), df.total_amount.cast("float"))
df_distance.write.insertInto("tripdb.distance")
```

Resultados en Hive:

```
hive> select * from payments limit 10;
OK
1      2021-01-01      2      11.8
1      2021-01-01      2      4.3
1      2021-01-01      1      51.95
1      2021-01-01      1      36.35
2      2021-01-01      1      24.36
1      2021-01-01      1      14.15
1      2021-01-01      2      17.3
1      2021-01-01      2      21.8
1      2021-01-01      4      28.8
1      2021-01-01      1      18.95
```

```
hive> select * from passenger limit 10 ;
OK
2021-01-01      1      11.8
2021-01-01      1      4.3
2021-01-01      1      51.95
2021-01-01      0      36.35
2021-01-01      1      24.36
2021-01-01      1      14.15
2021-01-01      1      17.3
2021-01-01      1      21.8
2021-01-01      1      28.8
2021-01-01      2      18.95
```

```
hive> select * from tolls limit 10;
OK
2021-01-01      1      0.0      11.8
2021-01-01      1      0.0      4.3
2021-01-01      1      0.0      51.95
2021-01-01      0      0.0      36.35
2021-01-01      1      0.0      24.36
2021-01-01      1      0.0      14.15
2021-01-01      1      0.0      17.3
2021-01-01      1      0.0      21.8
2021-01-01      1      0.0      28.8
2021-01-01      2      0.0      18.95
```

```
hive> select * from congestion limit 10;
OK
2021-01-01      1      2.5      11.8
2021-01-01      1      0.0      4.3
2021-01-01      1      0.0      51.95
2021-01-01      0      0.0      36.35
2021-01-01      1      2.5      24.36
2021-01-01      1      2.5      14.15
2021-01-01      1      0.0      17.3
2021-01-01      1      2.5      21.8
2021-01-01      1      0.0      28.8
2021-01-01      2      2.5      18.95
```

```
hive> select * from distance limit 10;
OK
2021-01-01      1      2.1      11.8
2021-01-01      1      0.2      4.3
2021-01-01      1      14.7      51.95
2021-01-01      0      10.6      36.35
2021-01-01      1      4.94      24.36
2021-01-01      1      1.6      14.15
2021-01-01      1      4.1      17.3
2021-01-01      1      5.7      21.8
2021-01-01      1      9.1      28.8
2021-01-01      2      2.7      18.95
```

5. Inserción de tablas mediante Pyspark con SQL:

Dropeamos la database

```
hive> drop database tripdb cascade;
OK
Time taken: 1.157 seconds
hive> show databases
> ;
OK
default
Time taken: 0.07 seconds, Fetched: 1 row(s)
```

Ahora se crea nuevamente con los comandos de Hive que se escribieron al principio.

```
hive> show tables;
OK
congestion
distance
passenger
payments
tolls
Time taken: 0.059 seconds, Fetched: 5 row(s)
```

Código para crear las tablas con SQL:

```
# Carga inicial del documento en hdfs a un dataframe llamado df
df = spark.read.option("header", "true").csv("/ingest/yellow_tripdata_2021-01.csv")
df.createOrReplaceTempView("vs_df")

# Carga de la tabla payments
df_payments = spark.sql("select cast(VendorID as int), cast(tpep_pickup_datetime as date),
cast(payment_type as int), cast(total_amount as float) from vs_df")
df_payments.createOrReplaceTempView("vs_payments_load")
spark.sql("insert into tripdb.payments select * from vs_payments_load")

# Carga de la tabla passenger
df_passengers = spark.sql("select cast(tpep_pickup_datetime as date), cast(passenger_count as int),
cast(total_amount as float) from vs_df")
df_passengers.createOrReplaceTempView("vs_passengers_load")
spark.sql("insert into tripdb.passenger select * from vs_passengers_load")

# Carga de la tabla tolls
df_tolls = spark.sql("select cast(tpep_pickup_datetime as date), cast(passenger_count as int),
cast(tolls_amount as float), cast(total_amount as float) from vs_df")
df_tolls.createOrReplaceTempView("vs_tolls_load")
spark.sql("insert into tripdb.tolls select * from vs_tolls_load")

# Carga de la tabla congestion
df_congestion = spark.sql("select cast(tpep_pickup_datetime as date), cast(passenger_count as int),
cast(congestion_surcharge as float), cast(total_amount as float) from vs_df")
df_congestion.createOrReplaceTempView("vs_congestion_load")
spark.sql("insert into tripdb.congestion select * from vs_congestion_load")

# Carga de la tabla distance
df_distance = spark.sql("select cast(tpep_pickup_datetime as date), cast(passenger_count as int),
cast(trip_distance as float), cast(total_amount as float) from vs_df")
df_distance.createOrReplaceTempView("vs_distance_load")
spark.sql("insert into tripdb.distance select* from vs_distance_load")
```

Los resultados en Hive son los mismos que llenando las tablas sin sentencias SQL:

```
hive> select * from payments limit 10;
OK
1      2021-01-01      2      11.8
1      2021-01-01      2      4.3
1      2021-01-01      1      51.95
1      2021-01-01      1      36.35
2      2021-01-01      1      24.36
1      2021-01-01      1      14.15
1      2021-01-01      2      17.3
1      2021-01-01      2      21.8
1      2021-01-01      4      28.8
1      2021-01-01      1      18.95
```

```
hive> select * from passenger limit 10;
OK
2021-01-01      1      11.8
2021-01-01      1      4.3
2021-01-01      1      51.95
2021-01-01      0      36.35
2021-01-01      1      24.36
2021-01-01      1      14.15
2021-01-01      1      17.3
2021-01-01      1      21.8
2021-01-01      1      28.8
2021-01-01      2      18.95
```

```
hive> select * from tolls limit 10;
OK
2021-01-01      1      0.0      11.8
2021-01-01      1      0.0      4.3
2021-01-01      1      0.0      51.95
2021-01-01      0      0.0      36.35
2021-01-01      1      0.0      24.36
2021-01-01      1      0.0      14.15
2021-01-01      1      0.0      17.3
2021-01-01      1      0.0      21.8
2021-01-01      1      0.0      28.8
2021-01-01      2      0.0      18.95
```

```
hive> select * from congestion limit 10;
OK
2021-01-01      1      2.5      11.8
2021-01-01      1      0.0      4.3
2021-01-01      1      0.0      51.95
2021-01-01      0      0.0      36.35
2021-01-01      1      2.5      24.36
2021-01-01      1      2.5      14.15
2021-01-01      1      0.0      17.3
2021-01-01      1      2.5      21.8
2021-01-01      1      0.0      28.8
2021-01-01      2      2.5      18.95
```

```
hive> select * from distance limit 10;
OK
2021-01-01      1      2.1      11.8
2021-01-01      1      0.2      4.3
2021-01-01      1      14.7      51.95
2021-01-01      0      10.6      36.35
2021-01-01      1      4.94      24.36
2021-01-01      1      1.6      14.15
2021-01-01      1      4.1      17.3
2021-01-01      1      5.7      21.8
2021-01-01      1      9.1      28.8
2021-01-01      2      2.7      18.95
```


Práctica de transformaciones con Spark y SQL

1. En el container de Nifi crear un .sh que permita descargar el archivo yellow_tripdata_2021-01.parquet.

```
nifi@c8a33460bcef:~/ingest$ cat ingest_trip_data.sh
wget https://dataengineerpublic.blob.core.windows.net/data-engineer/yellow_tripdata_2021-01.parquet
```

2. Por medio de Nifi crear un job con dos procesos:

- GetFile para obtener el archivo del punto 1.
- PutHDFS para ingesarlo a HDFS.



Se puede ver en el contenedor de hadoop hdfs que se cargó el .parquet correctamente.

```
hadoop@401bec58e4c6:/$ hdfs dfs -ls /nifi
Found 2 items
-rw-r--r-- 1 nifi supergroup 5462 2024-04-24 10:54 /nifi/starwars.csv
-rw-r--r-- 1 nifi supergroup 21686067 2024-05-08 08:51 /nifi/yellow_tripdata_2021-01.parquet
```

3. Con el archivo ingestado en HDFS/nifi, escribir consultas y agregar captura de pantalla del resultado. Para los ejercicios se puede usar SQL mediante la creación de una vista llamada yellow_tripdata.

```
df2 = spark.read.parquet("/nifi/yellow_tripdata_2021-01.parquet")
df2.createOrReplaceTempView("yellow_tripdata")
```

3.1 Mostrar las columnas VendorID int, Tpep_pickup_datetime date, Total_amount double, donde el total_amount < 10 USD.

```
df_3_1 = spark.sql("select cast(VendorID as int), cast(tpep_pickup_datetime as date), cast(total_amount as float) from yellow_tripdata where total_amount < 10")
```

```
>>> df_3_1.show(10)
+-----+-----+-----+
|VendorID|tpep_pickup_datetime|total_amount|
+-----+-----+-----+
|1|2020-12-31|4.3|
|2|2020-12-31|8.3|
|2|2020-12-31|9.96|
|2|2020-12-31|9.3|
|2|2020-12-31|5.8|
|1|2020-12-31|0.0|
|1|2020-12-31|9.3|
|2|2020-12-31|9.8|
|2|2020-12-31|8.8|
|2|2020-12-31|9.96|
+-----+-----+-----+
```

3.2 Mostrar los 10 días que más se recaudó dinero (tpep_pickup_datetime, total_amount).

```
df_3_2_cast = spark.sql("select cast(tpep_pickup_datetime as date), cast(total_amount as float) from yellow_tripdata")
df_3_2_cast.createOrReplaceTempView("3_2_2_cast")
df_3_2 = spark.sql("select tpep_pickup_datetime, sum(total_amount) as sum_total_amount from 3_2_2_cast group by tpep_pickup_datetime order by sum_total_amount desc")
```

```
>>> df_3_2.show(10)
+-----+-----+
|tpep_pickup_datetime| sum_total_amount|
+-----+-----+
|2021-01-28|961322.5573227406|
|2021-01-22|942205.9273384213|
|2021-01-29|937373.5074597001|
|2021-01-21|932444.4470283389|
|2021-01-15|931628.1872401834|
|2021-01-14|926664.0374201536|
|2021-01-27|895259.8676985204|
|2021-01-19|890581.4475597739|
|2021-01-07|887670.1573269963|
|2021-01-08|878002.7276113033|
+-----+-----+
```

3.3 Mostrar los 10 viajes que menos dinero recaudó en viajes mayores a 10 millas (trip_distance, total amount)

```
df_3_3_cast = spark.sql("select cast(trip_distance as float), cast(total_amount as float) from yellow_tripdata")
df_3_3_cast.createOrReplaceTempView("3_2_3_cast")
df_3_3 = spark.sql("select trip_distance, total_amount from 3_2_3_cast where trip_distance > 10 order by total_amount asc")
```

```
>>> df_3_3.show(10)
+-----+-----+
|trip_distance|total_amount|
+-----+-----+
|12.68| -252.3|
|34.35| -176.42|
|14.75| -152.8|
|33.96| -127.92|
|29.1| -119.3|
|26.94| -111.3|
|20.08| -107.8|
|19.55| -102.8|
|19.16| -90.55|
|25.83| -88.54|
+-----+-----+
```

3.4 Mostrar los viajes de más de dos pajeros que hayan pagado con tarjeta de crédito. (trip_distance, tpep_pickup_datetime)


```
df_3_4_cast = spark.sql("select cast(trip_distance as float), cast(tpep_pickup_datetime as date),  
cast(passenger_count as int), cast(payment_type as int) from yellow_tripdata")  
df_3_4_cast.createOrReplaceTempView("3_2_4_cast")  
df_3_4 = spark.sql("select trip_distance, tpep_pickup_datetime from 3_2_4_cast where  
passenger_count > 2 and payment_type = 1")
```

```
>>> df_3_4.show(10)  
+-----+-----+  
|trip_distance|tpep_pickup_datetime|  
+-----+-----+  
|          6.11|          2020-12-31|  
|           1.7|          2020-12-31|  
|          3.15|          2020-12-31|  
|         10.74|          2020-12-31|  
|           2.01|          2020-12-31|  
|           2.85|          2020-12-31|  
|           1.68|          2020-12-31|  
|           0.77|          2020-12-31|  
|            0.4|          2020-12-31|  
|         16.54|          2020-12-31|  
+-----+-----+
```

3.5 Mostrar los 7 viajes con mayor propina en distancias mayores a 10 millas
(tpep_pickup_datetime, trip_distance, passenger_count, tip_amount)

```
df_3_5 = spark.sql("select cast(tpep_pickup_datetime as date), cast(trip_distance as float),  
cast(passenger_count as int), cast (tip_amount as float) from yellow_tripdata where trip_distance >  
10 order by tip_amount DESC LIMIT 7")
```

```
>>> df_3_5.show()  
+-----+-----+-----+-----+  
|tpep_pickup_datetime|trip_distance|passenger_count|tip_amount|  
+-----+-----+-----+-----+  
|          2021-01-20|          427.7|             1|    1140.44|  
|          2021-01-03|          267.7|             1|     369.4|  
|          2021-01-12|          326.1|             0|     192.61|  
|          2021-01-19|          260.5|             1|     149.03|  
|          2021-01-31|           11.1|             0|     100.0|  
|          2021-01-01|          14.86|             2|      99.0|  
|          2021-01-18|           13.0|             0|      90.0|  
+-----+-----+-----+-----+
```

3.6 Mostrar para cada uno de los valores RateCodeID, el monto total y el monto promedio. Excluir los viajes donde RateCodeID es 'Group ride'

```
df_3_6 = spark.sql(" SELECT CAST(RateCodeID as int), SUM(total_amount),  
AVG(total_amount) FROM yellow_tripdata WHERE RateCodeID != 6 GROUP BY RateCodeID")
```

```
>>> df_3_6.show()  
+-----+-----+-----+  
|RateCodeID| sum(total_amount)| avg(total_amount)|  
+-----+-----+-----+  
|          1|1.9496468430212937E7|15.606626116946773|  
|          4|  90039.930000000082|  74.90842762063296|  
|          3|  67363.260000000043|  78.69539719626219|  
|          2|  973635.47000000732|  65.52937609369182|  
|         99| 1748.0699999999997|  48.55749999999999|  
|          5| 255075.089999999086|  48.939963545662096|  
+-----+-----+-----+
```